

# Vulnray - X-ray for code vulnerabilities

## Local AI scanning for foundational and low-level code

- Focused on C code: Bitcoin Core, hardware wallet firmware
- Runs locally with `VulnLLM-R-7B` reasoning model
- Takes under ~6 GB VRAM/shared memory
- Allows long-running repo-scale autonomous scans
- NOT burning tons of tokens
- NOT disclosing your findings to AI providers

## The Problem

- Security review is slow, expensive, and hard to parallelize across large repos
- Cloud LLM scanning is brittle for security work: privacy, cost, rate limits, and censorship risk
- Low-level C/C++ codebases hide subtle bugs: memory safety, integer bounds, crypto misuse

## Principles: what "Good" Looks Like

- Local-first CLI (automation, scripting, background execution)
- High-recall scanning (overnight / multi-day acceptable)
- Modes: `max-recall` , `balanced` , `deterministic` (CI-friendly, `temp=0`)
- Model-agnostic GGUF swap, `llama.cpp` backend (Metal on Apple Silicon)
- Outputs in one run: `json` , `csv` , `md` (with executive summary + findings table)

# How Vulnray Works

## End-to-end pipeline

```
CLI + config -> file discovery -> chunking (function/sliding) -> prompt  
profile + focus ->  
llama.cpp inference (VulnLLM-R-7B GGUF) -> parse findings -> reports  
(JSON/CSV/MD)
```

## Per-chunk contract

- Prompts request **JSON-only** outputs (machine-parseable findings + reasoning)

## Config precedence

```
CLI > ENV > TOML > defaults
```

## Running It (Tutorial Flow)

```
python -m pip install -e .
vulnray tutorial/test_project --config tutorial/vulnray.toml
```

Produces:

- tutorial/reports/demo\_scan.json
- tutorial/reports/demo\_scan.csv
- tutorial/reports/demo\_scan.md
- tutorial/reports/demo\_scan.prompt\_output.md (prompt + model output log)

# What You Get (Reports)

## Markdown report structure

- Executive summary (findings by severity)
- Findings table (triage-friendly)
- Per-finding detail: file, line range, function, CWE, severity, confidence, recommendation

## JSON report structure (scan metadata + findings)

- `scan_metadata` : tool, model, mode, timestamp, repo\_root, files scanned, chunks analyzed
- `findings[]` :  
`{id,file,start_line,end_line,function,vulnerability_type,severity,confidence,description,reasoning,recommendation,references}`

## Demo Scan Results (tutorial/test\_project)

From `tutorial/reports/demo_scan.md` (mode: `balanced`):

- **3 high-severity findings**
- CWE-787 out-of-bounds write risk (`strcpy` into fixed-size buffer)
- CWE-787 overflow risk in path construction (`sprintf` into fixed-size buffer)
- CWE-190 integer overflow risk (unchecked multiplication)

# Why VulnLLM-R (arXiv:2512.07533v1)

## Model claim (paper)

- "First specialized reasoning LLM" for vulnerability detection
- 7B-parameter reasoning model trained via a custom recipe (data selection, teacher-generated reasoning, filtering/correction, efficiency tuning)
- Teacher models (paper): DeepSeek-R1 and QwQ-32B
- Agent scaffold + context retrieval; paper reports outperforming traditional tools (e.g., CodeQL, AFL++) in real-world projects and discovering 15 zero-days

## Credit (authors)

Yuzhou Nie, Hongwei Li, Chengquan Guo, Ruizhe Jiang, Zhun Wang, Bo Li, Dawn Song, Wenbo Guo.